MiniScrub Documentation

Version 0.1.0
By: Nathan LaPierre
August 28, 2017

Table of Contents

Overview and Rationale

MiniScrub is a method for *de novo* scrubbing low-quality portions of long sequencing reads with deep learning. New sequencing technologies such as Oxford Nanopore MinION [1] and PacBio SMRT [2] have made high throughput long read sequencing feasible. These longer read lengths (1kbp to >100kbp) are better than short reads for analysis tasks such as assembly, alignment, and clustering due in part to their ability to span longer repeats [3]. However, these technologies are currently affected by high error rates, which can range for example from 5% to 40% per read [4].

MiniScrub provides several benefits:
- Entirely *de novo* (does not require a reference database) and thus can be used on unknown or ambiguous sequences
- Does not require an additional set of short reads, unlike the common hybrid error correction method
- Leverages Deep Learning, a powerful class of machine learning algorithms that have achieved record-breaking results in many fields such as computer vision and natural language processes, and is increasingly being applied in life sciences. Convolutional Neural Networks (CNNs), which we use in this project, have shown particularly promising results. See for example:
  - DeepVariant [5], a record breaking variant calling method using CNNs
  - Classifying skin lesion types from images with dermatologist-level accuracy using CNNs [6]
  - Using CNNs to detect mitosis in breast cancer histology images [7]

MiniScrub is meant to predict the percent identity of segments of reads, defined as the number of correct bases divided by the number of correct bases plus incorrect bases plus inserted and deleted bases, and trim or "scrub" the segments below a user defined threshold. First, minimap is used to find reads that have minimizers in common. Then, pileup images are generated using the minimap output and the reads file. Finally, a pre-trained Keras deep learning model is loaded and applied. The user sets a percent identity cutoff threshold (say, 0.75), and MiniScrub writes the reads file again without read segments that are predicted to be below the threshold.

MiniScrub is most effective when there is medium to high sequencing depth, and when cutting off very low quality read segments (i.e. >20% error rate). Running it will require a fairly significant amount of compute time, storage, and memory, and is thus best suited for use on cloud, cluster, grid, or supercomputer platforms that have these resources.

Installation and Dependencies

The MiniScrub files can be downloaded from GitHub: https://github.com/nlapier2/MiniScrub

If you are using git and would like to clone the repository, navigate to your command line (Mac, Linux, or Windows git client) and run the following line:

```
git clone https://github.com/nlapier2/MiniScrub.git
```

To download as a zip, use this link: https://github.com/nlapier2/MiniScrub/archive/master.zip
Then you can unzip it using either a file explorer window or on the command line using the "unzip" command (Mac or Linux).

MiniScrub is written in Python 2 and Python 3 and has the following dependencies:
- Standard python data science libraries: numpy, scipy, pandas, and scikit-learn
- Minimap (https://github.com/lh3/minimap)
- Keras (https://github.com/fchollet/keras)
- TensorFlow (https://github.com/tensorflow/tensorflow)
- If you are training the network (as opposed to just using an existing model), you must have SAM files with alignment information for the ground truth

If you are on the NERSC Cori supercomputer, you can load these dependencies by running the following:

```
module load python
module load deeplearning
```

If you already have a python model loaded, you may need to unload it first:

```
module unload python
```

Note that pileup.py uses Python 3 instead of Python 2. On Cori you would instead have to run:

```
module load python/3.6-anaconda-4.4
```

Basic Usage

Here we cover a basic sample usage of MiniScrub. Consult the "Major Scripts and Usage" section for a full list of options. Please note that these steps will take a very long time (up to several days) to run if you use the default settings and have many reads (i.e. over a million).

As a prerequisite to running MiniScrub, you must have a fastq file containing your reads. If you are also training the network, you must have a SAM file mapping those reads to reference sequences, so that this ground truth information can be used to train the network.

The lines below represent a typical run through of the MiniScrub pipeline. Please note that the step of training a network is optional, as a pretrained model will be made available online soon. We use the example names "reads.fastq" and "mapped-reads.sam" for the prerequisite files.

Running our modified version of minimap (generates PAF file with minimizer locations listed):
```
Temp
```

Generating pileup images and saving them in plots/ folder (please note that this script uses Python 3; on some computers, run python command instead of python3 command):
```
python3 pileup.py --mapping minimapped-reads.paf --reads
reads.fastq --saveplots --plotdir plots/
```

Loading a model and using it to scrub all 100bp segments with predicted 25% or higher error rate (--cutoff 0.75) using pileup images in plots/ folder:
```
python miniscrub.py --reads reads.fastq --cutoff 0.75
--load model.hd5 --input plots/ --output
scrubbed-reads.fastq
```

If training the network, execute these lines after generating the pileup images.

Generating labels for each 100bp segment of each read using SAM file:
```
python cigarToPctId.py --sam mapped-reads.sam --paf
minimapped-reads.paf --outfile labels.txt
```

Training a network on all images in plots/ folder:
```
python vgg.py --input plots/ --labels labels.txt --output
model.hd5
```

Major Scripts and Usage

Below are descriptions and help outputs for the major scripts that make up MiniScrub. Help outputs can be generated by running "`python [scriptname].py -h`". Many of the options (debug, limit_*, etc) have to do with speeding up the runtime by using only some reads, and are helpful for doing test runs. Other options provide for the ability to run a naive SVM baseline, based on minimizer counts, to test whether a newly-trained network is more accurate.

cigarToPctId.py

Description: This script is used to extract labels from SAM files that map the reads to reference sequences for training purposes. It reads the cigar strings and uses them to compute percent identities for read segments of user-defined size (must be at least 48, due to the constraints of the neural network). Additionally, these labels can be calculated in terms of an amount of bases or an amount of minimizers. You will only need this script if you are training a new model.

usage: cigarToPctId.py [-h] [--amount AMOUNT] [--compression {none,gzip}]
             [--limit_paf LIMIT_PAF] [--limit_sam LIMIT_SAM]
             [--mode {minimizers,bases}] [--outfile OUTFILE]
             [--paf PAF] --sam SAM

Use cigar strings in sam files to compute percent identities.

optional arguments:
 -h, --help   show this help message and exit
 --amount AMOUNT   Number of bases/minimizers per label.
 --compression {none,gzip}
                Compression format, or none
 --limit_paf LIMIT_PAF
                Optionally limit the number of reads from paf file.
 --limit_sam LIMIT_SAM
                Optionally limit the number of reads from sam file.
 --mode {minimizers,bases}
                Labels for minimizers or bases.
 --outfile OUTFILE   File to output label information to.
 --paf PAF   Path to paf file with minimizers.
 --sam SAM   Path to the sam file with ground truth labels.

<u>miniscrub.py</u>

Description: This is the script that performs the actual scrubbing of the reads. In addition to the basic usage previously listed, this script can also accept a labels file to load and test a model. This can be useful for seeing how well the method works on your dataset at various cutoffs when you have some labelled data available. The segment_size and window_size arguments are used if a new model is trained with different settings using cigarToPctId.py and vgg.py.

usage: miniscrub.py [-h] [--compression {none,gzip}] [--cutoff CUTOFF]
           [--debug DEBUG] [--input INPUT] [--labels LABELS] --load
           LOAD [--output OUTPUT] [--reads READS]
           [--segment_size SEGMENT_SIZE] [--window_size WINDOW_SIZE]

Use saved keras model to scrub reads.

optional arguments:
 -h, --help   show this help message and exit
 --compression {none,gzip}
              Compression of reads file ("none" or "gzip").
 --cutoff CUTOFF   Scrub read segments below this percent identity.
 --debug DEBUG   Number of images to use in debug mode. If <= 0, non-
           debug mode (default).
 --input INPUT   Directory with png pileup images. Default: current directory.
 --labels LABELS   Path to image labels file. If provided, will NOT trim.
           Labels must correspond with segment_size.
 --load LOAD   Path to keras model file to load. Required.
 --output OUTPUT   File to write scrubbed reads to.
 --reads READS   Path to reads file. Default: do not trim (output
           statsitics instead).
 --segment_size SEGMENT_SIZE
              Neural net segment size to predict. Keep as default
              unless network retrained.
 --window_size WINDOW_SIZE
              Neural net window size to predict. Keep as default
              unless network retrained.

pileup.py

Description: (NOTE: this script uses Python 3) This script generates pileup images of reads minimapped to each other (1 per read) that are used for training or applying a model. The --color can be set to bw (black and white) for faster generation and lower storage requirements. Conversely, setting the --maxdepth higher (must be at least 48) will result in slower generation and more storage space taken, but may slightly improve results.

   The --debug argument runs a short test run, and the --limit_* arguments can be used to limit the amount of reads scanned from the fastq and/or paf files. The --processes argument controls the degree of multiprocessing, which can lead to a large speedup, and it is highly recommended to set this as high as possible.

usage: pileup.py [-h] [--avgdepth AVGDEPTH] [--color {bw,rgb}]
   [--compression {none,gzip}] [--debug] [-k K]
   [--limit_fastq LIMIT_FASTQ] [--limit_reads LIMIT_READS]
   --mapping MAPPING [--maxdepth MAXDEPTH]
   [--mode {whole,minimizers}] [--plotdir PLOTDIR]
   [--processes PROCESSES] --reads READS [--saveplots] [--verbose]

Create pileups from .paf read-to-read mapping and fastq reads.

optional arguments:
 -h, --help   show this help message and exit
 --avgdepth AVGDEPTH Average coverage depth (not currently used).
 --color {bw,rgb} Color mode (bw or rgb).
 --compression {none,gzip}   Compression format, or none
 --debug   Debug mode.
 -k K   K-mer size of minimizers. Required.
 --limit_fastq LIMIT_FASTQ Limit number of reads to scan from fastq file.
 --limit_reads LIMIT_READS Limit number of reads to generate pileups for.
 --mapping MAPPING Path to the .paf file of read-to-read mappings.
 --maxdepth MAXDEPTH Maximum number of matched reads per pileup image.
 --mode {whole,minimizers}   Whole read or minimizers-only.
 --plotdir PLOTDIR If --saveplots is used, directory path to save plots in.
 --processes PROCESSES
    Number of multiple processes to run concurrently.
 --reads READS Path to the .fastq reads file.
 --saveplots If used, will plot the pileups and save them.
 --verbose Verbose output option.

<u>vgg.py</u>

Description: This script is used to train the CNN on pileup images. The CNN in question has the same architecture as VGG16 [8], except with an extra layer with a single ReLu in order to get a single real value as output. However, a --classify option is available if one desires to train the network for classification instead, in which case the last layer will be a single sigmoid instead.

      The user can specify the number of --epochs to train for, and has the option of adding --extra fully connected layers to the network, which can potentially lead to improved performance, but will slow training. Users can use the --load, --test_labels, and --test_input arguments to load an existing Keras model and test it on the specified data. The --baseline, --loadsvm, and --outputsvm arguments allow a user to additionally train, load, or save a naive SVM method for comparison with the neural network's results.

      The --segment_size option should be the same as the --amount option used in cigarToPctId.py, but the --window_size argument allows you the option of giving a larger "window" of the read than the actual segment to predict, which will slow down training but can allow the network to use nearby information to improve results. By default, these are set to 100 and 200, respectively, meaning that we give the network a 200bp image, the middle 100bp of which is what we are predicting percent identity for.

usage: vgg.py [-h] [--baseline] [--classify CLASSIFY] [--debug DEBUG]
        [--epochs EPOCHS] [--extra EXTRA] [--input INPUT]
        [--labels LABELS] [--load LOAD] [--loadsvm LOADSVM]
        [--output OUTPUT] [--outputsvm OUTPUTSVM]
        [--segment_size SEGMENT_SIZE] [--test_input TEST_INPUT]
        [--test_labels TEST_LABELS] [--window_size WINDOW_SIZE]

Create pileups from .paf read-to-read mapping and fastq reads.

optional arguments:
 -h, --help   show this help message and exit
 --baseline   Compare with baseline: SVM trained on # of minimizers matched.
 --classify CLASSIFY Turn into classification problem; specify threshold.
 --debug DEBUG   Number of examples to use in debug mode. If 0, non-
          debug mode (default).
 --epochs EPOCHS   Number of epochs to train the network. Default: 5.
 --extra EXTRA   Number of fully connected layers to add to VGG.
          Default: 0
 --input INPUT   Directory with png pileup images. Default: current directory.
 --labels LABELS   Path to image labels file.

--load LOAD   Path to keras model file to load. Default: do not load.

--loadsvm LOADSVM   Path to saved pickled SVM file. Default: do not load.

--output OUTPUT   File to write model to. Default: no output.

--outputsvm OUTPUTSVM

             File to write svm model to. Default: no output.

--segment_size SEGMENT_SIZE

             Size of read segments to evaluate.

--test_input TEST_INPUT

             Directory of serparate set of images to test model on.

--test_labels TEST_LABELS

             Path to file with labels for images in test set.

--window_size WINDOW_SIZE

             Window size for VGG. Window >= segment size.

References

1. https://nanoporetech.com/products/minion
2. http://www.pacb.com/smrt-science/
3. https://genomebiology.biomedcentral.com/articles/10.1186/gb-2013-14-6-405
4. http://genome.cshlp.org/content/25/11/1750.full
5. http://www.biorxiv.org/content/early/2016/12/21/092890
6. https://www.nature.com/nature/journal/v542/n7639/abs/nature21056.html
7. https://link.springer.com/chapter/10.1007/978-3-642-40763-5_51
8. https://arxiv.org/abs/1409.1556